

SUMMARY

Data Science graduate student at Rutgers University with 2+ years of industry experience and strong academic training in data analysis, machine learning, SQL, and analytics-driven application development. Experienced in building end-to-end data solutions, predictive models, and interactive dashboards.

EDUCATION

Rutgers University, New Brunswick, NJ, USA
Master of Science, Data Science | GPA: 3.83/4

2024 - E. May 2026

Manipal University Jaipur, Rajasthan, India
Bachelor of Technology in Information Technology | GPA: 8.07/10

2017- July 2021

PROFESSIONAL EXPERIENCE

Accenture, Gurugram, India

Analyst, Custom Software Engineering

May 2021 - July 2023

- Supported enterprise analytics platforms (Tableau, Cognos), improving reporting reliability and performance.
- Analyzed and resolved 100+ client-reported issues using ServiceNow by debugging data pipelines and application workflows.
- Monitored issue lifecycle metrics to ensure SLA compliance and improve client satisfaction.
- Performed system maintenance and upgrades, contributing to improved platform stability.

SKILLS

Languages & Querying: Python, SQL, R

Data Analysis & Visualization: Pandas, Exploratory Data Analysis (EDA), Data Cleaning, Feature Engineering

Databases & Platforms: ServiceNow, InstallShield, Visual Studio, Streamlit, GIT, PostgreSQL, MySQL, Snowflake

Analytics Concepts: Predictive Modeling, Data Wrangling, Causal Inference, Data Visualization, ETL, LLM

ACADEMIC PROJECTS

Carbon Emissions Tracking in Language Model Fine-Tuning for Question Answering Task | Python, Hugging Face, CodeCarbon | [\[To Repo\]](#)

- Evaluated performance vs. carbon emissions trade-offs for transformer models on SQuAD v2.0 QA tasks.
- Fine-tuned DistilBERT, BERT, GPT and RoBERTa using full fine-tuning, LoRA, and few-shot learning.
- Integrated CodeCarbon to track GPU/CPU/RAM energy usage and CO₂ emissions during training.
- Achieved ~95% of full fine-tuning performance with 12–24% lower emissions using LoRA.
- Reported efficiency metrics including F1 score per kg CO₂, highlighting sustainable ML practices.

Extracting Financial Sentiment Features with Sparse Autoencoders | [\[To Repo\]](#)

- Applied Sparse Autoencoders (SAEs) to GPT-2 internal activations to extract interpretable financial sentiment features.
- Addressed polysemy and superposition by enforcing sparsity constraints on latent representations.
- Extracted and analyzed residual stream activations from GPT-2 layer 9 using the FinancialPhraseBank dataset.
- Evaluated feature interpretability using firing rate analysis, top-activating tokens, and domain-specific financial scoring.
- Demonstrated the feasibility of extracting domain-relevant, interpretable features from black-box language models.

Anime & Manga Recommendation and Community Platform | Python, Streamlit, PostgreSQL, SQLAlchemy | [\[To Repo\]](#)

- Built a data-driven web application for browsing, reviewing, and recommending anime and manga content.
- Designed a relational PostgreSQL schema for users, content, reviews, and activity tracking with optimized queries.
- Implemented search, filtering, sorting, and pagination to improve performance and usability.
- Developed a rule-based shuffle recommender and an OpenAI-powered natural language search assistant that converts user queries into structured SQL filters for personalized recommendations.

PROFESSIONAL CERTIFICATES

snowflake

SnowPro Core Certification (Aug 2025 - Aug 2027) - <https://achieve.snowflake.com/f93aa466-1e48-4236-a287-fbfb880b7dd8#acc.TD4Uz15Z>